

1 Solution Strategy

For this project, I implemented a Conflict-Driven Clause Learning (CDCL) SAT solver in Python, later optimized with Cython. I then added some specific augmentations to generic CDCL drawing inspiration from the MiniSAT solver [Sorensson and Een \(2005\)](#); [Sörensson and Claessen \(2026\)](#). After reading a survey paper on restart strategies [Haim and Heule \(2014\)](#), I spent (perhaps too much) time thinking about different restart ideas. Most of these did not pan out. Eventually, I ended up with an implementation with the following core components:

- I used Variable State Independent Decaying Sum (VSIDS) for choosing branching variables. I learned about this starting at a set of lecture slides on SAT Heuristics [Clarke](#), from which I learned about the Chaff solver [Moskewicz et al. \(2001\)](#). My VSIDS implementation is not as sophisticated as the Chaff solver, but has the general flavor.
- Conflict analysis: I used a First Unique Implication Point (1-UIP) scheme, that I learned about from a set of lecture slides from the CIS189 course at the University of Pennsylvania [Menezes \(2024\)](#).
- Restart strategies: Luby sequence-based restarts to escape poor search regions [Luby et al. \(1993\)](#); [Haim and Heule \(2014\)](#).

2 What Worked

My biggest wins in terms of speedup were:

Activity and Clause-Quality Heuristics (VSIDS + LBD) I use the VSIDS heuristic to prioritize variables that appear in recent conflicts, biasing the search toward recently active parts of the formula. Complementing this, I also implemented a simple version of the Literal Block Distance metric [Chang et al. \(2017\)](#) to evaluate learned clauses by the number of distinct decision levels they span.

The Luby restart sequence which provides short initial probes that grow geometrically. This helped balance exploration of the search space and exploitation of the value of lots of fast initial restarts [Haim and Heule \(2014\)](#).

Negative Branching Direction Following MiniSat’s approach, I always assign the decision variables to false. Haim et al [Haim and Heule \(2014\)](#) show this is better than arbitrary, and that false branching empirically outperforms true branching. This simple heuristic provided noticeable speedups versus random or true heuristics.

Low-Level Representational Changes After compiling the Python source to C via Cython, I implemented several representational changes to optimize the C-level execution. This primarily involved replacing ”boxed” Python structures (e.g., `dict`, `Optional`) with raw C-style constructs.

Specifically, the variable assignment mapping was transitioned from a dictionary-based structure, assignments: Dict[int, Optional[bool]], to a contiguous integer array, assignments: List[int]. In this representation, an assignment for a variable v is defined as 0 if unassigned, 1 if True, and -1 if False.

3 What Did Not Work

I spent *most* of my time trying to tune instance-adaptive hyperparameters based on clause-to-variable density. This approach was inspired by portfolio solvers like SATzilla [Xu et al. \(2008\)](#), which select different solvers based on instance features. My adaptation applies the same principle within a single solver by switching hyperparameters at startup as well as learning about how clause density could affect the probability of a formula's SAT or UNSAT-ness [Dudek et al. \(2017\)](#).

Initially, I attempted to use a phase transition heuristic based on the known 3-SAT phase transition at density 4.26 [Dudek et al. \(2017\)](#). I hypothesized that this could guide restart strategies or other parameters. However, this didn't help in practice—the harder instances in the test suite weren't hitting this specific density point. I eventually ended up classifying instances into three broad clusters based on their clause density, and assigned roughly the following hyperparams:

- Density ≤ 6 : Likely SAT, Slow variable decay (0.99), large Luby base (512), and low randomization (0.005).
- Density (6–30): Probably in the phase-transition region and hard to solve (decay 0.95, Luby 100, random 0.02)
- High density (> 30): Likely UNSAT. Aggressive decay (0.85), frequent restarts (Luby 32), and high randomization (0.08)

Unfortunately, most of the harder CNF files I tested tended to be in the medium density / phase transition range. High and low density instances are already effectively dealt with by clause learning. There's no real reason to optimize for them. At the same time, I do feel like there were modest gains from this approach.

Giving UNSAT for Extreme Density I briefly considered returning UNSAT immediately for extremely high-density instances, trading completeness for speed. While this would be correct with high probability and extremely fast on some UNSAT cases, I ultimately rejected this approach because it would make the solver incomplete.

4 How Much Time Did I Spend ?

I've implemented DPLL and CDCL solvers before for other classes, so an initial implementation only took me about an hour. After this, however, I spend several hours trying to speed it up and/or experiment with phase transitions. I estimate I spent about 10 hours on this project.

References

- Wenjing Chang, Guanfeng Wu, and Yang Xu. Adding a lbd-based rewarding mechanism in branching heuristic for sat solvers. In *2017 12th International Conference on Intelligent Systems and Knowledge Engineering (ISKE)*, pages 1–6, 2017. doi: 10.1109/ISKE.2017.8258780.
- Edmund M. Clarke. Heuristics for efficient SAT solving. Lecture slides, Carnegie Mellon University. URL https://www.cs.cmu.edu/~emc/15-820A/reading/sat_cmu.pdf. Accessed: 2026-02-15.
- Jeffrey M. Dudek, Kuldeep S. Meel, and Moshe Y. Vardi. Combining the k -cnf and xor phase-transitions, 2017. URL <https://arxiv.org/abs/1702.08392>.
- Shai Haim and Marijn Heule. Towards ultra rapid restarts. *arXiv preprint arXiv:1402.4413*, 2014.
- Michael Luby, Alistair Sinclair, and David Zuckerman. Optimal speedup of las vegas algorithms. *Information Processing Letters*, 47(4):173–180, 1993.
- Rohan Menezes. Lecture 4: Modern techniques in SAT solving. University of Pennsylvania, CIS 1890 Course Materials, 2024. URL <https://www.cis.upenn.edu/~cis1890/files/Lecture4.pdf>. Accessed: February 15, 2026.
- Matthew W Moskewicz, Conor F Madigan, Ying Zhao, Lintao Zhang, and Sharad Malik. Chaff: Engineering an efficient sat solver. In *Proceedings of the 38th annual Design Automation Conference*, pages 530–535, 2001.
- Niklas Sörensson and Koen Claessen. MiniSat: A minimalistic and high-performance SAT solver. <https://github.com/niklasso/minisat>, 2026. Accessed: 2026-02-15.
- Niklas Sorensson and Niklas Een. Minisat v1. 13-a sat solver with conflict-clause minimization. *SAT*, 53(2005):1–2, 2005.
- Lin Xu, Frank Hutter, Holger H Hoos, and Kevin Leyton-Brown. Satzilla: portfolio-based algorithm selection for sat. *Journal of artificial intelligence research*, 32:565–606, 2008.